

Papillon

AI-Powered Cybersecurity Platform

Final Project Report

Course: BIL-496

Date: 17 April 2026

Institution: TOBB University of Economics and Technology

Department: Department of Computer Engineering

Team Members:

- Kerem Elma - 211101079
- Mehmet Eski - 221104095
- Mete Cem Turan - 211101054
- Sabri Mert Pişkin - 211101035

Semester: Spring 2025-2026

Github Repository: <https://github.com/SabriMertPiskin/Papillon>

Table of Contents

1. Introduction	4
1.1 Project Overview	4
1.2 Project Objectives	4
1.3 Scope	4
1.4 Definitions, Acronyms, and Abbreviations	5
2. Final Architecture and Design	5
2.1 System Architecture	5
2.2 Backend Architecture (Server)	6
2.2.1 SSH Based AWS VM Lab	6
2.2.2 Local VM Lab	7
2.2.3 Attack Surface Analysis Module	8
2.2.4 CVE Management Module	11
2.2.5 Blacklist Management Module	12
2.2.6 Cryptography Module	14
2.2.7 Outlook Integration Module	15
2.3 Frontend Architecture (Client)	17
2.4 AI/ML Module Architecture	17
2.4.1 Malware Detection	17
2.4.2 Password Strength Module	20
2.4.3 Phishing Detection Module	23
2.5 Database Schema	26
2.6 API Endpoint Inventory	26
2.7 Routing and Navigation	27
3. Final Project Status	28
3.1 Module Status Summary	28
3.2 Feature Completion Matrix	28
3.3 Codebase Statistics	29
4. Impact of Engineering Solutions	29
4.1 Global Context	29
4.2 Economic Context	29
4.3 Environmental Context	29
4.4 Societal Context	30
5. Contemporary Issues	30
5.1 Rise of AI-Powered Cyber Threats	30
5.2 Zero-Day Vulnerability Economy	30
5.3 Regulatory Compliance and Data Privacy	30
5.4 Supply Chain Security	30
6. Tools and Technologies	31
6.1 Backend Technologies	31
6.2 Frontend Technologies	31
6.3 AI/ML Technologies	31
6.4 DevOps and Infrastructure	31

6.5 Security Libraries	31
7. Background Research and References	32
7.1 Library Resources	32
7.2 Internet Resources and APIs	32
7.3 Similar Systems and Design Inspiration	32
7.4 Engineering Principles Applied	32
8. Test Results	33
8.1 Test Environment	33
8.2 Test Summary Statistics	33
8.3 Detailed Test Results by Module	33
8.3.1 Users — Authentication, RBAC, MFA (FR10 TC-04, TC-12)	33
8.3.2 Password AI — Strength Assessment (FR11 TC-11)	36
8.3.3 Attack Surface Analysis (FR3 TC-10)	36
8.3.4 Cryptographic Services (FR8 TC-09)	36
8.3.5 Malware Analysis (FR7 TC-03)	37
8.3.6 IP Blacklist Management (FR1 TC-08)	37
8.3.7 Network Intrusion Detection (FR4 TC-01)	37
8.3.8 Phishing Detection (FR6 TC-02)	37
8.3.9 CVE Vulnerability Tracking (FR2 TC-06)	37
8.4.3 Discussion	38
9. Conclusion	39
9.1 Achievements	39
9.2 Known Limitations	39
9.3 Future Enhancements	39
10. User's Manual	40
10.1 Login Page	40
10.2 Dashboard	40
10.3 Password Analysis	41
10.4 Text Encrypter	41
10.5 CVE	42
10.6 Vulnerability Map	43
10.7 Attack Surface	43
10.8 cPanel Data	44
10.9 VM Lab	44
10.10 SSH VM Lab	45
10.11 Outlook Integration	46
10.12 Phishing History	46
10.13 Malware Analysis	47
10.14 Profile & Account	48
10.15 Blacklist	48
11. References	49
11.1 Datasets	50

1. Introduction

1.1 Project Overview

Papillon is an AI-powered cybersecurity platform designed to provide comprehensive security monitoring, threat detection, and vulnerability management through a unified web-based dashboard. The platform integrates three distinct machine learning models with traditional security tools, enabling real-time threat analysis across multiple attack vectors including network intrusions, malware, phishing emails, and weak credentials. Offering highly flexible deployment and connectivity options, the platform allows users to effortlessly set up and connect to an AWS Virtual Machine (VM) or connect directly to local VMs.

1.2 Project Objectives

- **Unified Security Dashboard** — Consolidate multiple cybersecurity functions into a single, intuitive web interface.
- **AI-Driven Threat Detection** — Integrate machine learning models for automated classification of network intrusions, malware, phishing attempts, and password vulnerability.
- **Real-Time Monitoring** — Provide live analysis capabilities for network traffic, CVE vulnerabilities, and attack surface reconnaissance.
- **Role-Based Access Control** — Implement RBAC to segregate administrative and analyst functionalities.
- **Multi-Factor Authentication** — Enforce TOTP-based MFA to secure user accounts.
- **Extensible Architecture** — Design a modular system that can accommodate new security modules without architectural changes.
- **Flexible Deployment & Integration:** Enable seamless setup and direct connectivity with AWS Virtual Machines and local VM environments to support versatile hosting and monitoring strategies.

1.3 Scope

The Papillon platform encompasses the following functional domains:

- **Authentication & Authorization:** User registration, login, MFA (TOTP), RBAC (Admin/Analyst), session management
- **Vulnerability Intelligence:** Real-time CVE fetching from NVD API, vulnerability mapping across network topology
- **Cryptographic Services:** AES-256-GCM encryption/decryption, RSA-2048, SHA-256/512, MD5, Base64
- **Reconnaissance:** Attack surface scanning (DNS, SSL, subdomains, open ports, WHOIS)
- **AI Threat Detection:** cPanel, Malware (RandomForest), Phishing Detection (TF-IDF + Random Forest via ONNX), Password Strength (RandomForest)
- **Network Management:** IP Blacklist CRUD, cPanel integration, VM Lab management
- **Communication:** Microsoft Outlook OAuth integration for email threat monitoring

1.4 Definitions, Acronyms, and Abbreviations

Term	Definition
RBAC	Role-Based Access Control
MFA	Multi-Factor Authentication
TOTP	Time-based One-Time Password
NVD	National Vulnerability Database

CVE	Common Vulnerabilities and Exposures
CVSS	Common Vulnerability Scoring System
IDS	Intrusion Detection System
ONNX	Open Neural Network Exchange
REST	Representational State Transfer
AES-GCM	Advanced Encryption Standard — Galois/Counter Mode
RSA	Rivest-Shamir-Adleman cryptosystem
CIDR	Classless Inter-Domain Routing
OAuth	Open Authorization protocol
SPA	Single Page Application

2. Final Architecture and Design

2.1 System Architecture

Papillon follows a three-tier architecture consisting of a React-based frontend and a Django REST backend, and a MySQL database with external AI model files.




Figure 1: Papillon three-tier system architecture showing the Presentation, Application, and Data tiers.

2.2 Backend Architecture (Server)

The Django backend is organized into 11 independent Django applications, each responsible for a specific security domain:

Django App	Key Dependencies	Responsibility
users/	pyotp, qrcode	Authentication, session management, MFA, RBAC, profile
cve/	requests	Fetching and normalizing CVE data from NVD REST API
crypto/	cryptography	AES-256-GCM, RSA-2048, hashing (SHA, MD5, Base64)
attack_surface/	dnspython, bs4	DNS, subdomain enumeration, SSL scanning, port scanning, etc.
phishing/	onnxruntime	AI-based phishing email classification and history
password_ai/	scikit-learn	Random Forest based password strength assessment
malware_analysis/	scikit-learn	Random Forest based malware detection
network_ids/	requests, numpy	cPanel integration network monitoring
blacklist/	requests	IP blacklist CRUD with IPv4/IPv6/CIDR validation
outlook/	requests	Microsoft Outlook OAuth integration
vm_lab/	boto3	Virtual machine lab management via AWS API
ssh_vm_lab/	django.http , boto3, pathlib/os	RBAC-protected AWS EC2 start trigger and secure .pem SSH key download endpoint.

Design Patterns Applied:

- **Lazy Loading:** All AI models are loaded on first request and cached in memory to avoid startup overhead.
- **Graceful Degradation:** If external tools or AI models fail, views return structured error responses (HTTP 503) instead of crashing.
- **CSRF Exemption:** API endpoints use session-based auth with CORS instead of CSRF tokens for compatibility.

2.2.1 SSH Based AWS VM Lab

2.2.1.1 Overview

This project is a management module developed using the Django web framework that allows users to remotely start EC2 virtual machine instances on Amazon Web Services (AWS). It also provides a secure method for downloading the necessary AWS key files (.pem) required for SSH access. The system is designed to meet the needs of cybersecurity laboratories or remote VM Lab environments.

2.2.1.2 Architecture

2.2.1.2.1. URL Configuration (urls.py)

The system exposes two primary endpoints to the users:

- **start-instance/:** Triggers the process to power on the virtual machine.

- **download-key/**: Serves the .pem AWS key file required for the SSH connection.

2.2.1.2.2. View Layer and Security (views.py)

The view layer handles the core business logic and security protocols:

- **Role-Based Access Control (RBAC)**: Using the `@require_role('admin','analyst')` decorator, critical operations are restricted only to users with 'admin' or 'analyst' privileges.
- **Instance Activation**: The `start_instance_view` function retrieves the `AWS_INSTANCE_ID` from environment variables and invokes the `start_ec2_instance` service. The result is returned to the user as a JSON response.
- **Secure File Serving**: The `download_aws_key` function verifies the existence of the .pem file on the server using `AWS_KEY_PATH.exists()`. If the file is present, it is securely served as a downloadable attachment using `FileResponse`.

2.2.1.3 Operational Workflow

1. An authorized user logs into the web interface.
2. The user clicks the "Start Instance" button, which activates the Ubuntu server on AWS.
3. The user downloads the `awskey.pem` file through the provided link.
4. The user connects to the remote server via their local terminal using the command:

```
ssh -i "awskey.pem" ec2-user@<aws-ip>.
```

2.2.2 Local VM Lab

2.2.2.1 Overview

This report details the implementation of a local agent architecture designed to bridge the gap between a centralized Django management system and local virtualization software (VirtualBox).

2.2.2.2 Architecture Overview

The system operates on a Manager-Agent architecture:

- **The Manager (Django)**: Acts as the orchestrator, handling user authentication, role-based access (RBAC), and session state management.
- **The Agent (Flask)**: A lightweight Python script running on the host machine where the VMs are located. It listens for requests and executes system-level commands to start or stop virtual machines.

2.2.2.3 Detailed Component Analysis

2.2.2.3.1 Local Agent Implementation (agent.py)

The agent is built using Flask and utilizes the subprocess module to interface with the VBoxManage CLI tool.

- **Endpoint /start-vm:** Receives the path for the VirtualBox executable and initiates the 'startvm' command with the --type gui flag.
- **Endpoint /stop-vm:** Executes the 'controlvm' poweroff command to safely terminate the virtual machine session.

2.2.2.4 Operational Workflow

1. **Agent Activation:** The agent.py script is started on the host machine.
2. **User Request:** An authorized user calls the VM Lab start endpoint.
3. **Authentication & Validation:** Django verifies the user's role and checks if a session already exists.
4. **Agent Trigger:** Django sends an HTTP request to the local Flask agent.
5. **Subprocess Execution:** The agent runs the VBoxManage command, launching the "django-kali" VM.
6. **State Update:** Django updates the internal session tracker and returns the VM connection details to the user.

2.2.3 Attack Surface Analysis Module

The Attack Surface Analysis Module is responsible for identifying externally reachable assets and exposure points of a target domain. Its purpose is to support security assessment by collecting a consolidated view of DNS configuration, subdomains, SSL/TLS status, open ports, email exposure, administrative interfaces, robots.txt content, and basic IP information. Rather than relying on a monolithic scanning class, the implementation is organized as a lightweight orchestration layer in the Django backend that delegates specialized tasks to separate utility scripts. This separation improves maintainability, testability, and clarity of responsibilities.

The main orchestration logic is implemented in the Django application layer. The attack_surface_scan endpoint authenticates the user, determines the target domain according to the user role, resolves the domain to an IP address, and then invokes the underlying analysis utilities one by one. The domain selection logic is handled through helper functions that normalize input values, validate whether a requested target is allowed, and apply role-based fallbacks. Analysts can scan their own configured domain, while administrators must explicitly specify the analyst account they are acting on behalf of. This design enforces access control at the scanning entry point and reduces the risk of unauthorized target selection.

The auxiliary analysis utilities are stored in the attack_surface_analysis_tools directory. Each file focuses on a single security task. The DNS utility resolves common record types such as A, AAAA, NS, CNAME, MX, TXT, SOA, and CAA records. The subdomain utility performs parallel checks against a predefined subdomain list to identify live hosts. The SSL scanner validates certificate information and derives a basic certificate validity result. The port scanner checks a small set of commonly exposed ports and uses either nmap or socket-based fallbacks

depending on availability. The email harvester extracts publicly visible email addresses from the target website and its common informational pages. The admin panel scanner probes well-known administrative paths and collects response metadata that may indicate exposed login interfaces. The robots parser retrieves the robots.txt file, the Whois module extracts registration information, and the IP utility resolves the target domain to an IPv4 address.

From a software engineering perspective, the module is intentionally functional rather than class-heavy. Most components are implemented as standalone functions so that each security task can be executed independently and tested in isolation. This makes the code easier to extend: if a new reconnaissance step is needed, it can be added as a separate utility and called from the orchestration layer without affecting the rest of the system. The current structure also matches the endpoint's runtime behavior, where results from multiple lightweight utilities are aggregated into a single JSON response.

The module is covered by unit and integration tests focused on domain normalization, access control, and endpoint behavior. These tests verify that domain formatting is handled consistently, that foreign targets are rejected, and that the scan endpoint behaves correctly for authenticated analysts and administrators. As a result, the Attack Surface Analysis Module acts as a controlled security reconnaissance layer within the backend architecture.

If you want, I can also turn this into a full subsection set with headings and short paragraphs for each utility so you can paste it directly into the report section by section.

2.2.3.1 DNS Record Resolution

This utility is responsible for identifying the DNS configuration of the target domain. It resolves common record types, including A, AAAA, NS, CNAME, MX, TXT, SOA, and CAA, in order to provide a broad view of how the domain is configured on the public internet. From a security perspective, this information helps reveal the hosting structure, mail routing, verification records, and other metadata that may be useful during reconnaissance. The implementation is intentionally lightweight and function-based, which makes the DNS lookup process easy to isolate, test, and reuse within the larger attack surface workflow.

2.2.3.2 Subdomain Enumeration

This module is used to discover active subdomains belonging to the target domain. It compares a predefined list of candidate subdomains against the target and checks which hosts respond successfully. The implementation performs these checks in parallel to improve efficiency and reduce scan time. In an academic context, this component contributes to the discovery of additional exposed services and infrastructure entry points that may not be visible from the primary domain alone. It is therefore a key part of the reconnaissance phase of the overall attack surface analysis.

2.2.3.3 SSL/TLS Certificate Analysis

This utility inspects the SSL/TLS certificate presented by the target domain and evaluates whether the certificate is currently valid. The scanner collects certificate-related metadata and checks the expiration status, which helps determine whether the domain is protected by a properly maintained encrypted transport layer. This is important because outdated or misconfigured certificates may indicate poor operational hygiene or a weakened security

posture. The module provides both structured scan output and a simple validity result, allowing the backend to present certificate findings in a clear and usable format.

2.2.3.4 Port Scanning

This component identifies which commonly used network ports are open on the target IP address. It focuses on a predefined set of ports that are frequently associated with web services, remote administration, databases, and file-sharing protocols. If the nmap library is available, the scanner uses it for richer results; otherwise, it falls back to socket-based connection checks. This design ensures that the scan remains functional even in limited environments. From a security standpoint, open ports are one of the most direct indicators of exposed services and potential attack entry points.

2.2.3.5 Email Harvesting

This module extracts publicly visible email addresses from the target website. It scans the homepage and a small set of common information pages such as contact, about, and support pages, then filters and normalizes the discovered addresses. It also detects mailto links in the HTML content. In the context of attack surface analysis, this feature helps identify exposed contact information that may be relevant for phishing risk assessment, social engineering exposure, or organizational footprint mapping. The implementation is intentionally limited to a small set of likely pages to keep the scan non-intrusive and efficient.

2.2.3.6 Admin Panel Scanning

This utility probes a set of common administrative paths such as admin, login, wp-admin, dashboard, cms, and similar routes. It records response metadata such as HTTP status, redirects, and header/body hints that may indicate the presence of a login page or management interface. This does not attempt exploitation; instead, it serves as a reconnaissance tool for identifying publicly reachable administrative surfaces. In the report, this module can be described as a lightweight exposure detector for administrative endpoints.

2.2.3.7 Robots.txt Retrieval

This module fetches the target domain's robots.txt file and returns its contents. Although simple, this information can still be useful because robots files often disclose restricted directories, crawler hints, or application paths that are not immediately visible through normal browsing. As part of attack surface analysis, it provides another source of passive reconnaissance data without performing intrusive actions.

2.2.3.8 Whois Lookup

This component retrieves domain registration information through a Whois query. It exposes metadata such as registration dates, registrar details, and other publicly available domain ownership fields. From a security analysis perspective, Whois data can help establish domain age, ownership context, and lifecycle information that may be relevant when evaluating trustworthiness or operational maturity. The module converts the raw Whois output into a structured list so that it can be displayed consistently in the final scan results.

2.2.3.9 IP Resolution

This utility resolves the domain to its IPv4 address. The returned IP address is used as a foundation for other tasks, especially port scanning and general infrastructure mapping. Although technically simple, it is an important prerequisite for the rest of the attack surface workflow because many external checks are performed at the network level rather than at the DNS name level. The module therefore acts as the bridge between domain-based reconnaissance and host-based analysis.

2.2.4 CVE Management Module

The CVE Management Module is responsible for retrieving, normalizing, and presenting recent vulnerability information from the National Vulnerability Database (NVD). Its primary goal is to provide security analysts with up-to-date vulnerability intelligence in a structured and readable format. Unlike modules that store data locally, this component acts as a live integration layer that queries an external vulnerability source on demand and transforms the raw response into a simplified JSON structure. This design keeps the module lightweight while ensuring that the information remains current.

The architecture of the module is intentionally simple and service-oriented. A single Django view function serves as the public API entry point, while the NVD REST API acts as the external data source. The endpoint is exposed without authentication, which makes it suitable for general vulnerability awareness and dashboard-style reporting. The view handles request validation, result filtering, data normalization, and response formatting. In this way, the module separates API communication concerns from presentation concerns and keeps the implementation easy to maintain.

Detailed Component Analysis

2.2.4.1 latest_cves Endpoint

The latest_cves endpoint is the core implementation of the module. It accepts a GET request and allows the caller to specify a result limit through a query parameter. The limit is constrained to a safe range to prevent excessive output and unnecessary API load. The endpoint then queries the NVD API for CVE entries published within a recent time window and iterates through paginated results until enough valid vulnerabilities have been collected.

During this process, the endpoint performs several normalization steps. Rejected or deferred CVEs are excluded, English descriptions are preferred when multiple languages are available, and severity values are derived in a priority order from CVSS v4.0, CVSS v3.1, and CVSS v2 fields. The endpoint also collects publication dates, external references, and the NVD detail URL for each CVE. The final output is returned as a structured JSON object containing the count, the applied limit, and the normalized CVE list.

From a software design perspective, the endpoint is function-based rather than class-based. This makes the module straightforward to test and keeps the implementation focused on data retrieval and transformation. If the project later requires persistent storage, caching, or scheduled synchronization, those responsibilities can be introduced as additional services without changing the public response contract of the endpoint.

2.2.4.2 Operational Workflow

The workflow begins when a user or frontend component sends a request to the CVE endpoint. The backend reads the optional limit parameter, prepares the NVD query window, and fetches vulnerability records from the external API. The returned results are filtered and normalized into a simplified format. After this, the backend constructs a final JSON response containing the most relevant fields for reporting and analysis. This workflow ensures that raw NVD data is converted into application-friendly output before being exposed to the user interface.

2.2.4.3 Validation and Testing

The module is covered by tests that validate the public endpoint, response normalization, severity parsing, language selection, limit handling, pagination behavior, and error handling. These tests confirm that the endpoint returns consistent data even when the external API response changes in structure or content. As a result, the CVE Management Module provides a reliable vulnerability intelligence layer within the overall backend architecture.

2.2.5 Blacklist Management Module

The Blacklist Management Module is responsible for managing blocked IP addresses within the backend system. Its main purpose is to allow authenticated users to maintain a structured list of malicious or suspicious IP addresses and to enrich this list with external reputation data when needed. The module supports both local blacklist operations and an integrated AbuseIPDB lookup workflow, making it a practical security control layer for incident response and threat handling.

The architecture of the module is intentionally lightweight and service-oriented. A Django view layer exposes the public API, while a dedicated database model stores blacklist entries and preserves their metadata. The implementation follows a functional design rather than a class-heavy design, which keeps the code simple, testable, and easy to extend. Each responsibility is separated into small endpoints: one for listing and creating blacklist entries, one for querying AbuseIPDB, and one for deleting entries. This separation improves maintainability and makes the module easier to reason about in a security-oriented application.

Detailed Component Analysis

2.2.5.1 Blacklist Storage Model

The BlacklistedIP model is the persistence layer of the module. It stores the blocked IP address, the blocking reason, the username of the user who created the entry, and the timestamp of creation. The `ip_address` field is unique, which prevents duplicate blacklist records and ensures that each blocked host is represented only once. The model accepts IPv4, IPv6, and CIDR-style addresses, allowing the blacklist to represent both individual hosts and address ranges. Since the model is kept minimal, it focuses only on data storage and does not introduce unnecessary business logic.

2.2.5.2 Blacklist List/Create Endpoint

The `blacklist_list_create` endpoint handles both reading and writing operations. On GET requests, it returns the full blacklist as a JSON list, including the stored metadata for each entry. On POST requests, it validates the incoming IP address, checks whether the format is acceptable, prevents duplicates, and then creates a new blacklist entry. If no custom reason is provided, the module assigns a default reason such as “Manual block.” This endpoint is protected by session-based authentication, so only authenticated users can access it.

2.2.5.3 AbuseIPDB Lookup Endpoint

The `blacklist_abuseipdb_lookup` endpoint provides external reputation enrichment. Instead of storing data locally, it sends the supplied IP address to the AbuseIPDB API and retrieves reputation-related information such as abuse confidence score, total reports, number of distinct reporters, country code, usage type, ISP, and last report time. This endpoint is useful when the system needs to assess whether an IP address has already been flagged by the wider security community. In the module architecture, this endpoint acts as an external intelligence layer rather than a storage mechanism.

2.2.5.4 Delete Endpoint

The `blacklist_delete` endpoint removes an existing blacklist entry by its database identifier. This operation allows analysts to maintain the list over time and remove entries that are no longer relevant. The endpoint also uses session authentication and returns clear success or error responses depending on whether the target entry exists. From a design perspective, this endpoint completes the module’s basic CRUD-like functionality, even though the blacklist is intentionally limited to the operations needed for operational security tasks.

2.2.5.5 Operational Workflow

The workflow begins when an authenticated user opens the blacklist interface or submits a new IP address for blocking. The backend first validates the session and then checks whether the provided address matches the accepted IPv4, IPv6, or CIDR patterns. If the input is valid and not already stored, the system saves it as a blacklist record. When external reputation data is needed, the user can invoke the AbuseIPDB lookup endpoint, which sends the IP address to the third-party service and returns a normalized response. If the user later decides to remove an entry, the delete endpoint deletes the corresponding record from the database. This workflow keeps blacklist management straightforward while still supporting both internal and external threat intelligence.

2.2.5.6 Validation and Testing

The module is covered by tests that verify IP address validation, authentication behavior, duplicate detection, list retrieval, record creation, delete operations, and response handling for invalid inputs. The tests also confirm that the module supports common IPv4 and IPv6 formats and rejects malformed values. This test coverage helps ensure that the blacklist remains consistent and that invalid or duplicated security records do not enter the system.

2.2.6 Cryptography Module

The Cryptography Module provides a controlled backend service for encrypting, decrypting, hashing, and encoding text data. Its purpose is to demonstrate and support common cryptographic operations within the platform while preserving per-user traceability through database storage. The module is implemented as a small Django-based service layer and exposes two main endpoints for encryption and decryption. Rather than relying on a complex object-oriented hierarchy, the implementation uses a functional design in which each cryptographic operation is handled directly by the view layer. This keeps the module lightweight, easy to test, and straightforward to extend.

The architecture of the module combines application-level request handling with a persistent storage model. The Django view layer receives the user request, validates the session, selects the requested algorithm, and executes the corresponding cryptographic routine. The EncryptedData model stores the operation metadata and output values so that encryption events can be tracked historically. This design allows the module to behave both as a cryptographic utility and as a record-keeping component for security analysis and demonstration purposes.

Detailed Component Analysis

2.2.6.1 EncryptedData Model

The EncryptedData model is the persistence layer of the module. It stores the user who performed the operation, the selected algorithm, the plaintext preview, and the resulting ciphertext or hash output. For algorithms such as AES and RSA, it also stores the generated key material required for later decryption. For one-way functions such as MD5, SHA-1, SHA-256, and SHA-512, it stores the hash result instead. The model is intentionally generic and supports multiple algorithm types through a single unified schema, which simplifies storage and retrieval.

2.2.6.2 Encrypt Endpoint

The `encrypt_text` endpoint is the main encryption entry point. It accepts a POST request containing the input text and the desired algorithm. The endpoint first validates authentication and verifies that plaintext was provided. It then performs one of several supported operations. For AES-256-GCM, it generates a random 256-bit key and nonce, encrypts the plaintext, and stores the ciphertext together with the key and nonce. For RSA-2048, it generates a key pair, encrypts the plaintext using the public key, and returns the corresponding public and private key material. For MD5, SHA-1, SHA-256, and SHA-512, it computes a hash digest. For Base64, it performs encoding rather than encryption. After each operation, the result is stored in the database and returned as a JSON response.

2.2.6.3 Decrypt Endpoint

The `decrypt_text` endpoint handles the inverse operation where applicable. It accepts a POST request containing the ciphertext and the corresponding algorithm. For AES-256-GCM, the endpoint requires the key and nonce in addition to the ciphertext. For RSA-2048, it requires the private key. For Base64, it performs decoding. The module explicitly rejects decryption

attempts for hash-based algorithms, since these are one-way functions and cannot be reversed. This clear separation between reversible and irreversible operations is important from a security design perspective because it prevents misuse and helps users understand the purpose of each cryptographic primitive.

2.2.6.4 Operational Workflow

The workflow begins when an authenticated user submits a plaintext value and chooses an algorithm. The backend validates the request, executes the selected cryptographic routine, stores the result in the database, and returns the output to the caller. If the user later submits a decryption request, the backend verifies that the algorithm supports reversal and then applies the required decryption or decoding process. This workflow makes the module suitable for both educational use and security-oriented demonstrations, while also preserving traceability of performed operations.

2.2.7 Outlook Integration Module

The Outlook Integration Module provides backend support for connecting a user account to Microsoft Outlook through the Microsoft Graph ecosystem. Its purpose is to enable OAuth-based account linking, encrypted token storage, connection status tracking, and mail retrieval from a connected mailbox. The module is implemented as a Django service layer that coordinates the OAuth flow, stores credentials securely, and exposes endpoints for managing the connection lifecycle. Instead of embedding the entire authentication flow inside a single function, the module separates credential persistence, authorization, callback handling, and mailbox access into distinct responsibilities.

The architecture follows a service-and-storage model. The Django view layer manages user interaction and communicates with Microsoft's OAuth and Graph endpoints through HTTP requests. The OutlookAccount model stores the encrypted access token, refresh token, and per-user OAuth configuration. Sensitive fields are protected using Fernet-based encryption, which ensures that credentials are not stored in plain text. This design makes the module appropriate for a security-sensitive application because it combines external identity integration with encrypted local persistence.

Detailed Component Analysis

2.2.7.1 OutlookAccount Model

The OutlookAccount model stores all Outlook-related state for a user. It keeps encrypted access and refresh tokens, encrypted client credentials, the account's expiration time, a connection flag, and the Outlook email address associated with the account. The model uses property-based accessors to transparently encrypt and decrypt values when they are written or read. This keeps encryption concerns centralized and prevents sensitive values from being exposed directly in application logic. The model is the core persistence component of the module and acts as the secure storage layer for the OAuth workflow.

2.2.7.2 Save Client Credentials

The `save_client_id` endpoint allows an authenticated analyst to submit a client ID and client secret before starting the OAuth flow. The values are temporarily stored in the session and, if an Outlook account already exists, they are also persisted in encrypted form through the model. This step prepares the system for later authorization and token exchange. In practical terms, it bridges the gap between user-provided OAuth configuration and the backend-managed Outlook connection lifecycle.

2.2.7.3 Authorization Flow

The `authorize` endpoint generates the Microsoft authorization URL and returns it to the caller. The endpoint verifies that the session contains the temporary client identifier, stores the current user context in the session, and constructs an OAuth authorization request with the required Graph scopes. This stage does not perform login itself; instead, it redirects the user toward Microsoft's authorization infrastructure. As a result, the module follows the standard OAuth pattern in which the backend prepares the authorization request and the identity provider handles user authentication.

2.2.7.4 OAuth Callback

The callback endpoint receives the authorization code after the user approves access on Microsoft's side. The backend exchanges the code for access and refresh tokens, retrieves basic mailbox identity data from Microsoft Graph, and stores the resulting connection state in the `OutlookAccount` model. At this stage, the account is marked as connected and the token expiration time is recorded. This callback completes the OAuth handshake and establishes a persistent link between the application user and the Outlook account.

2.2.7.5 Status and Mail Retrieval

The `get_outlook_status` endpoint reports whether an Outlook account is connected and returns the stored mailbox identity. The `disconnect` endpoint removes the association entirely, allowing the user to revoke the connection from the application side. The `get_latest_mail` endpoint retrieves the most recent email from the connected mailbox and refreshes the access token proactively if the stored token has expired. Together, these endpoints provide the operational lifecycle of the Outlook integration, from connection setup to status monitoring and mailbox access.

2.2.7.6 Operational Workflow

The workflow begins when an authenticated analyst submits Microsoft OAuth client credentials. The backend stores the credentials securely, generates an authorization URL, and directs the user to Microsoft for consent. After the user approves access, Microsoft returns an authorization code to the callback endpoint, which exchanges the code for tokens and persists the connection in encrypted form. Once connected, the system can report account status, refresh expired tokens automatically, retrieve the latest message, and disconnect the account when needed. This workflow enables secure and maintainable Outlook integration while keeping sensitive data encrypted at rest.

2.3 Frontend Architecture (Client)

The React frontend is a Single Page Application built with Vite, using React Router for client-side navigation. Key characteristics:

- 21 page components (Login, Register, Dashboard, CVE, Encryption, Attack Surface, Vulnerability Map, Network Traffic, Malware Analysis, Password Strength, Blacklist, Phishing History, Outlook, cPanel, VM Lab, User Profile, AsyncTaskExample, MFASettings, NotFound, OutlookCallback, SSHVMLab)
- 19 dedicated CSS files totaling ~140 KB
- Centralized API layer in services/api.js with Axios interceptors for auto-logout on 401
- Error Boundary component for graceful crash recovery
- RBAC-aware Sidebar that adapts navigation based on user role

2.4 AI/ML Module Architecture

Papillon integrates three pre-trained machine learning models, each serving a distinct security function:

Model	Algorithm	Format	Size	Input/Output
Malware Detection	Random Forest	.pkl	~500 MB	Executable file -> malware score
Phishing Detection	TF-IDF + Random Forest	.onnx	~5.3 MB	Email text -> phishing probability (0-1)
Password Strength	Random Forest	.pkl	~86 MB	Password -> strength level

2.4.1 Malware Detection

2.4.1.1 Overview

The Malware Analysis module performs static analysis on uploaded files using Random Forest and model trained on the EMBER dataset. It extracts a 512-dimensional feature vector from raw file bytes combining byte histogram and sliding window entropy to classify files as malicious, suspicious, or safe.

Property	Value
Algorithm	Random Forest (n_estimators=300, max_depth=40, min_samples_split=5, min_samples_leaf=2, max_features="sqrt", bootstrap=True, n_jobs=-1, random_state=42)
Format / Size	.pkl (joblib) / ~340 MB
Input	Executable file upload (multipart/form-data, max 100 MB)
Feature Vector	512-dim: 256-bin byte histogram + 256 entropy windows
Output	Verdict (malicious/suspicious/safe) + score (0-100)

2.4.1.2 Training Data

The model is trained on the EMBER dataset, a large-scale benchmark for static malware detection developed by Endgame (now Elastic).

Property	Value
Dataset	EMBER (Endgame Malware BEhavior Representation)
Total Samples	1,000,000+ PE (Portable Executable) files
Classes	Benign / Malicious (binary classification)
Feature Type	Static analysis (Byte histogram, Entropy)

2.4.1.3 Training Pipeline

2.4.1.3.1 Byte Histogram (Features 0-255)

A 256-bin histogram counting the frequency of each byte value (0x00-0xFF) in the file. This captures the overall byte distribution pattern, which differs significantly between compiled executables, packed/obfuscated malware, and normal documents.

```
hist = np.histogram(list(bytez), bins=256, range=(0, 255))[0]
```

2.4.1.3.2 Sliding Window Entropy (Features 256-511)

Shannon entropy is calculated over 1024-byte sliding windows across the file. The entropy values are padded/truncated to exactly 256 values. High-entropy regions indicate encryption, compression, or obfuscation common in malware. Low-entropy regions suggest plaintext or structured data.

```
window = 1024
byte_entropy = [shannon_entropy(bytez[i:i+window])
                 for i in range(0, len(bytez), window)]
byte_entropy = np.pad(byte_entropy, (0, max(0, 256 - len)))[0:256]
```

2.4.1.3.3 Entropy Calculation

Shannon entropy formula: $H = -\sum (p_i * \log_2(p_i))$ where p_i is the probability of byte value i . Returns 0.0 for uniform data (single byte), ~8.0 for random data (all 256 values equally distributed).

Data Type	Entropy Range	Interpretation
Uniform (all same byte)	0.0	No information content
Structured text (ASCII)	2.0 - 5.0	Normal human-readable data
Compiled code (.exe)	5.0 - 7.0	Typical executable or document
Encrypted / packed	7.0 - 8.0	Obfuscation or compression detected

Random data	~8.0	Maximum information density
-------------	------	-----------------------------

2.4.1.4 Inference Data Flow

End-to-end flow from file upload to response:

- **1. Upload:** User uploads executable file
- **2. Validation:** Auth check (@require_role('analyst')), file presence check, size limit (100 MB)
- **3. Reading:** Read file bytes
- **4. Feature Extraction:** Build 512-dim vector: 256-bin byte histogram + 256 sliding window entropy values
- **5. AI Prediction:** Lazy-loaded Random Forest model (.pkl) predicts class (0=benign, 1=malicious) + predict_proba for probability
- **6. Verdict:** Three-tier verdict based on prediction and probability: malicious (pred=1, prob>0.7), suspicious (pred=1 or prob>0.4), safe (otherwise). Score = probability * 100.
- **7. AI Tags:** Rule-based tags: High Confidence Malware, Obfuscated/Packed Code (entropy>7), Suspicious Small Executable (<50KB), Trusted File Structure, Normal Entropy

2.4.1.5 File Structure

File	Purpose
ember_model.pkl	Trained EMBER gradient boosted trees model (~340 MB)
ember_model.zip	Compressed model archive (~73 MB)
pipeline.py	Original training pipeline (feature extraction + prediction)
feature_extraction.ipynb	Jupyter notebook for feature engineering experiments
api.py	Standalone API wrapper (legacy)

Django File	Purpose
views.py	analyze_file(), _extract_features(), _byte_entropy(), _compute_hashes(), _generate_ai_tags()
tests.py	27 unit + integration tests
urls.py	/ai/malware/analyze/

2.4.1.6 Test Results

Class	Precision	Recall	F1-Score
Benign	0.914	0.893	0.903
Malicious	0.895	0.916	0.905
Weighted Avg.	0.905	0.904	0.904

Accuracy	0.904
ROC AUC	0.968

2.4.2 Password Strength Module

2.4.2.1 Overview

The Password Strength module uses RanveerChaudhary/password_strength_dataset from HuggingFace and assesses password security using a Random Forest classifier trained on password features. It extracts 10 numerical features from the input password and predicts a three-level strength classification: Weak (0), Normal (1), or Strong (2).

Property	Value
Algorithm	Random Forest Classifier
Format / Size	.pkl (joblib) / ~86 MB
Input	Password -> 10 engineered features from raw password
Output	Strength level: 0 (Weak), 1 (Normal), 2 (Strong)

2.4.2.2 Feature Engineering

The module extracts 10 features from the raw password string before passing them to the ML model:

#	Feature	Type	Description
1	length	int	Total character count

2	has_upper	bool (0/1)	Contains uppercase letter (A-Z)
3	has_lower	bool (0/1)	Contains lowercase letter (a-z)
4	has_digit	bool (0/1)	Contains digit (0-9)
5	has_special	bool (0/1)	Contains special char (!@#\$%^&*)
6	digit_count	int	Number of digits
7	special_count	int	Number of special characters
8	entropy	float	Shannon-based entropy (bits)
9	unique_chars	int	Count of distinct characters
10	repeated_ratio	float	length / unique_chars (lower = better)

Entropy Calculation

Password entropy is calculated based on the character set size and password length using Shannon entropy formula: $\text{entropy} = \text{length} * \log_2(\text{charset_size})$

Character Set	Charset Size	Example
Lowercase (a-z)	+26	password
Uppercase (A-Z)	+26	Password
Turkish lowercase (çğıöşu)	+6	sifre123
Turkish uppercase (CGIOSU)	+6	SIFRE123
Digits (0-9)	+10	pass123
Special (!@#\$%^&*...)	+32	pass@123

Example: 'StrongP@ss1' uses $\text{lowercase}(26) + \text{uppercase}(26) + \text{digit}(10) + \text{special}(32) = 94$ charset.
Entropy = $11 * \log_2(94) = 72.1$ bits.

2.4.2.3. Inference Data Flow

- **1. Request:** User types password in PasswordStrength.jsx -> POST /ai/password-strength/predict/

- **2. Validation:** JSON parsing, password extraction. Returns 400 if empty/missing. No auth required (public endpoint).
- **3. Feature Extraction:** `_extract_features()` computes 10-feature DataFrame from raw password string
- **4. Entropy:** `_calculate_entropy()` computes Shannon entropy based on detected charset
- **5. AI Prediction:** Lazy-loaded Random Forest model (.pkl) predicts strength level (0/1/2)
- **6. Suggestions:** `_generate_suggestions()` provides rule-based improvement feedback
- **7. Response:** Return JSON with prediction, strength_level, suggestions[], entropy, password_length

Example response:

```
{
  "success": true,
  "prediction": "Strong",
  "strength_level": 2,
  "suggestions": ["Great password! Consider enabling MFA for extra security."],
  "entropy": 72.1,
  "password_length": 11
}
```

2.4.2.4. Results

	Precision	Recall	F1
Weak	0.98	0.88	0.92
Normal	0.88	0.98	0.93
Strong	1	0.99	0.99

2.4.2.5. Suggestion Engine (Explainable AI)

The module generates educational feedback based on rule-based analysis of the password:

Condition	Suggestion
length < 8	Password should be at least 8 characters long.
8 <= length < 12	Consider using 12+ characters for stronger security.
No uppercase	Add uppercase letters (A-Z).
No lowercase	Add lowercase letters (a-z).
No digits	Add numbers (0-9).
No special chars	Add special characters (!@#\$\$%^&*).

unique_ratio < 0.5	Too many repeated characters. Use more variety.
In common list	This is a commonly used password. Choose something unique.
Strong + no issues	Great password! Consider enabling MFA for extra security.

2.4.2.6. File Structure

File	Purpose
password_strength_model.pkl	Trained Random Forest model (~86 MB)
pipeline.py	Original training pipeline (feature extraction + prediction)
api.py	Standalone API wrapper (legacy)

Django File	Purpose
views.py	predict_strength(), _calculate_entropy(), _extract_features(), _generate_suggestions()
tests.py	35 unit + integration tests
urls.py	/ai/password-strength/predict/

2.4.3 Phishing Detection Module

2.4.3.1. Overview

The Phishing Detection module classifies emails as phishing, suspicious, or clean using a TF-IDF + Random Forest pipeline trained on 164,906 emails, exported to ONNX for inference.

Property	Value
Algorithm	TF-IDF (bigram) + RandomForest (200 trees)
Format / Size	ONNX / ~5.3 MB
Accuracy	96.71%
Output	Probability scores [P(safe), P(phishing)]
Risk Tiers	Phishing (>=70) Suspicious (40-69) Clean (<40)

2.4.3.2. Training Datasets

7 public email datasets combined (~249 MB raw CSV):

Dataset	Size	Description
phishing_email.csv	101.7 MB	Large mixed phishing/legitimate corpus
CEAS_08.csv	64.8 MB	CEAS 2008 spam/phishing challenge
Enron.csv	43.5 MB	Enron email corpus (legitimate)
SpamAssassin.csv	14.2 MB	SpamAssassin public corpus
Ling.csv	8.9 MB	LingSpam dataset
Nigerian_Fraud.csv	8.8 MB	Nigerian 419 scam emails
Nazario.csv	7.5 MB	Jose Nazario phishing corpus

2.4.3.3. Training Pipeline

Preprocessing: Dynamic column detection per CSV -> concatenation -> dropna + strip + min length filter (>10 chars) -> label normalization (phishing/spam/1/yes -> 1, else -> 0) -> class balancing (undersample if ratio > 3:1) -> 80/20 stratified split (train: 131,924, test: 32,982).

2.4.3.3.1 Model Parameters

Stage	Parameter	Value
TF-IDF	max_features	8,000
TF-IDF	ngram_range	(1, 2) — unigrams + bigrams
TF-IDF	sublinear_tf	True — log normalization
TF-IDF	min_df / max_df	3 / 0.95
Random Forest	n_estimators	200 trees
Random Forest	max_depth	30
Random Forest	min_samples_leaf	2

2.4.3.3.2 ONNX Export

Pipeline exported with zipmap=False for raw probability arrays. ONNX outputs: [0] 'label' (0/1), [1] 'probabilities' [P(safe), P(phishing)].

2.4.3.4. Inference Data Flow

End-to-end flow from user input to response:

- **1. Request:** User submits email via PhishingHistory.jsx -> POST /ai/phishing/predict/
- **2. Validation:** Session auth check (@require_role('analyst')), JSON parsing, email_text validation
- **3. AI Inference:** PhishingDetector (Singleton) runs ONNX session -> returns label + [P(safe), P(phishing)]
- **4. Scoring:** score = round(p_phishing * 100), classify into phishing/suspicious/clean
- **5. XAI Reasons:** Rule-based enrichment: sender domain check, urgency keyword detection in subject
- **6. DB Save:** Save to PhishingLog (user, email_text, sender, subject, prediction, confidence, reasons)
- **7. Response:** Return JSON with status, score, confidence, p_phishing, p_safe, reasons[]

Example response:

```
{
  "status": "phishing", "score": 96,
  "confidence": 0.9599, "p_phishing": 0.9599, "p_safe": 0.0401,
  "reasons": ["Sender domain not recognized...", "AI model classified as phishing..."]
}
```

2.4.3.5. Performance Results

Class	Precision	Recall	F1-Score	Support
Safe	0.9751	0.9558	0.9654	15,834
Phishing	0.9599	0.9775	0.9686	17,148
Overall Accuracy			0.9671	32,982

2.4.3.5.1 Confusion Matrix

	Predicted Safe	Predicted Phishing
Actual Safe	TN = 15,134	FP = 700

Actual Phishing	FN = 386	TP = 16,762
-----------------	----------	-------------

False Positive Rate: 4.42% | False Negative Rate: 2.25% | True Positive Rate (Recall): 97.75%

2.5 Database Schema

The application uses the following database models:

- **CustomUser** (custom model): username, email, role (admin/analyst), domain, vm_lab_path, mfa_enabled, mfa_secret, mfa_backup_code
- **BlacklistedIP**: ip_address (unique, IPv4/IPv6/CIDR), reason, blocked_by, created_at
- **EncryptedData**: user (FK), key, nonce, private_key, public_key, hash_result
- **PhishingLog**: user (FK), sender, subject, preview, full_body, score, status, ai_label, ai_reasons, scanned_at
- **DomainTrafficEvent**: domain, client_ip, method, path, status_code, response_ms, user_agent, requested_at, source
- **CPanelCredential**: user (FK), host, username, token_encrypted, password_encrypted, verify_ssl, created_at, updated_at

2.6 API Endpoint Inventory

The backend exposes multiple REST API endpoints organized by domain.

Group	Endpoint	Method	Auth
Auth	/auth/register/	POST	No
	/auth/login/	POST	No
	/auth/logout/	POST	Yes
	/auth/dashboard/	GET	Yes
	/auth/change-password/	POST	Yes
	/auth/update-domain/	POST	Yes
	/auth/update-vm-lab-path/	POST	Yes
MFA	/auth/mfa/setup/	POST	Yes
	/auth/mfa/verify-setup/	POST	Yes
	/auth/mfa/verify/	POST	No
	/auth/mfa/disable/	POST	Yes
	/auth/mfa/status/	GET	Yes
CVE	/cve/latest/	GET	No
Crypto	/crypto/encrypt/	POST	Yes
	/crypto/decrypt/	POST	Yes
Attack Surface	/attack-surface/scan/	POST	Yes
AI	/ai/phishing/predict/	POST	Yes
	/ai/phishing/history/	GET	Yes
	/ai/password-strength/predict/	POST	No
	/ai/malware/analyze/	POST	Yes
	/ai/network-ids/predict/	POST	Yes
	/ai/network-ids/analyze-batch/	POST	Yes
Blacklist	/blacklist/	GET, POST	Yes
	/blacklist/<id>/	DELETE	Yes
	/blacklist/abuseipdb/	POST	Yes
VM Lab	/vm-lab/	GET, POST	Yes

	/vm-lab/status	GET	Yes
	/vm-lab/start	POST	Yes
	/vm-lab/terminate	POST	Yes
	/vm-lab/start-instance	GET	Yes
Outlook	/outlook/	GET, POST	Yes
	/outlook/save-client-id	POST	Yes
	/outlook/authorize	GET	Yes
	/outlook/callback	GET	Yes
	/outlook/disconnect	POST	Yes
	/outlook/status	GET	Yes
	/outlook/latest-mail	GET	Yes
SSH VM Lab	/ssh-vm-lab/	GET	Yes
	/ssh-vm-lab/start-instance/	GET	Yes
	/ssh-vm-lab/download-key/	GET	Yes

2.7 Routing and Navigation

The frontend uses React Router v6 with route guards:

Route	Component	Access
/login	Login	Public
/register	Register	Public
/dashboard	Dashboard	Protected
/cve	CVEList	Protected
/encryption	Encryption	Protected
/attack-surface	AttackSurfaceAnalysis	Protected
/vulnerability-map	VulnerabilityMap	Protected
/network-traffic	NetworkTraffic	Protected
/malware-analysis	MalwareAnalysis	Protected
/password-strength	PasswordStrength	Protected
/blacklist	Blacklist	Protected
/phishing-history	PhishingHistory	Protected
/outlook-integration	OutlookIntegration	Protected
/cpanel-data	CPanelData	Protected
/vm-lab	VMLab	Protected
/profile	UserProfile	Protected
/ssh-vm-lab	SSHVMLab	Protected
/async-demo	AsyncTaskExample	Protected
/outlook/callback	OutlookCallback	Protected
/mfa-settings	Navigate(to /profile)	Public redirect
/	Navigate(to dashboard or /login)	Public redirect
*	NotFound	Public (404)

3. Final Project Status

3.1 Module Status Summary

All modules have been implemented; module-level testing has been performed.

#	Module	Backend	Frontend	AI	Data Source	Status
1	Authentication	users/	Yes	—	Database	Complete
2	Dashboard	users/	Yes	—	Database	Complete
3	CVE Vulnerabilities	cve/	Yes	—	NVD API	Complete
4	Encryption	crypto/	Yes	—	Backend	Complete
5	Attack Surface	attack_surface/	Yes	—	Live DNS/WHOIS/SSL/Port scan data	Complete
6	Vulnerability Map	cve/	Yes	—	NVD -> topology	Complete
7	Malware Analysis	malware_analysis/	Yes	Random Forest	File upload	Complete
8	Password Strength	password_ai/	Yes	Random Forest	Input password	Complete
9	Phishing Detection	phishing/	Yes	ONNX	AI classify	Complete
10	IP Blacklist	blacklist/	Yes	—	DB CRUD	Complete
11	MFA Settings	users/	Yes	—	TOTP	Complete
12	Profile & Account	users/	Yes	—	Database	Complete
13	Outlook Integration	outlook/	Yes	—	MS OAuth	Complete
14	cPanel Data	network_ids/	Yes	—	cPanel API	Complete
15	VM Lab	vm_lab/	Yes	—	VirtualBox CLI	Complete
16	SSH VM Lab	ssh_vm_lab/	Yes	—	AWS API	Complete

Overall Completion: 100%

3.2 Feature Completion Matrix

ID	Functional Requirement	Status	Notes
FR1	IP Blacklist Management	Complete	CRUD with IPv4/IPv6/CIDR validation
FR2	CVE Vulnerability Tracking	Complete	Live NVD API integration with pagination
FR3	Attack Surface Reconnaissance	Complete	DNS, SSL, subdomain, port scanning
FR4	Network Intrusion Detection	Complete	cPanel integration
FR5	Virtual Machine Lab	Complete	AWS EC2 integration
FR6	Phishing Email Detection	Complete	ONNX-based phishing classification.
FR7	Malware File Analysis	Complete	EMBER-based malware model (.pkl).
FR8	Cryptographic Services	Complete	AES-256-GCM, RSA-2048, SHA, MD5
FR9	Outlook Email Integration	Complete	Microsoft OAuth 2.0 + Graph API (Mail.Read/User.Read).
FR10	Security & Access Control	Complete	RBAC (Admin/Analyst), session auth

FR11	Password Strength Assessment	Complete	ML-based password strength scoring with feature extraction.
------	------------------------------	----------	---

3.3 Codebase Statistics

Metric	Count
Backend Django Apps	12
Frontend Pages (JSX)	21
CSS Style Files	19
API Endpoints	46
AI Models Integrated	4
Database Models	7
Unit/Integration Tests	250+
Test Modules	10
Git Commits (main)	100+

4. Impact of Engineering Solutions

4.1 Global Context

The Papillon platform addresses a critical global need for accessible cybersecurity tools. According to the World Economic Forum Global Risks Report 2025 [1], cybercrime is projected to cost the world \$10.5 trillion annually, making it the third-largest "economy" after the US and China.

Papillon democratizes advanced cybersecurity capabilities by integrating AI-driven threat detection into a single, web-based platform that can be deployed in organizations regardless of geographical location. The use of open-source technologies (Django, React, scikit-learn) ensures that the platform is accessible without prohibitive licensing costs, particularly benefiting institutions in developing countries that lack dedicated cybersecurity infrastructure.

4.2 Economic Context

Traditional cybersecurity platforms (Splunk, CrowdStrike, Palo Alto Cortex) cost \$50,000-\$500,000+ annually in enterprise licensing. Papillon provides comparable functionality — IDS, malware analysis, phishing detection, vulnerability scanning — using open-source components, reducing the financial barrier to entry by over 90%.

By consolidating six or more separate security tools into one dashboard, Papillon reduces the context-switching overhead for security analysts. The AI-driven automation eliminates manual traffic inspection, reducing mean-time-to-detect (MTTD) from hours to seconds.

Small and medium enterprises, which constitute 43% of cyberattack targets but often cannot afford dedicated security teams, can leverage Papillon's automated AI modules to achieve baseline security monitoring without specialized personnel [2].

4.3 Environmental Context

Papillon's consolidated architecture means organizations deploy one platform instead of multiple separate security appliances, reducing server hardware requirements, power consumption, and electronic waste.

The lazy-loading strategy for AI models ensures CPU resources are consumed only when actively analyzing threats. These machine learning models are lightweight compared to deep learning alternatives, requiring no GPU acceleration and running efficiently on standard server hardware.

4.4 Societal Context

Papillon processes all data locally, email analysis, file scanning, and network monitoring occur on the organization's own infrastructure. No user data is transmitted to third-party cloud services for AI processing, preserving data sovereignty and privacy.

The Password Strength module provides educational feedback (entropy calculations, character class analysis, common password warnings), contributing to general cybersecurity awareness. By providing affordable IDS and vulnerability management, Papillon can be deployed in hospitals, schools, and municipal networks that are increasingly targeted by ransomware but lack dedicated security budgets.

5. Contemporary Issues

5.1 Rise of AI-Powered Cyber Threats

The same machine learning technologies that power Papillon's defenses are increasingly weaponized by threat actors. AI-generated phishing emails now bypass traditional keyword-based filters with 85%+ success rates. Papillon addresses this by using an ONNX-based TF-IDF + Random Forest model that evaluates semantic patterns rather than simple keyword matching, providing resilience against AI-crafted attacks.

5.2 Zero-Day Vulnerability Economy

The NVD publishes 60+ new CVEs daily, overwhelming security teams [3]. Papillon's Vulnerability Map automatically fetches, categorizes, and distributes CVEs across a visual network topology, enabling rapid prioritization based on CVSS severity scores. This addresses the "alert fatigue" problem that causes 55% of critical vulnerabilities to go unpatched for 30+ days.

5.3 Regulatory Compliance and Data Privacy

Regulations such as GDPR (EU), KVKK (Turkey), and CCPA (California) impose strict requirements on data protection. Papillon's on-premises architecture, AES-256-GCM encryption module, and MFA enforcement directly support compliance requirements. The attack surface scanning module helps organizations identify exposed data endpoints before regulators do.

5.4 Supply Chain Security

The SolarWinds and Log4Shell incidents demonstrated that vulnerabilities in third-party dependencies can cascade across entire ecosystems. Papillon's CVE module specifically tracks library-level vulnerabilities using data from NVD, and the Vulnerability Map's keyword-based distribution helps teams quickly identify which infrastructure components are affected by newly disclosed supply chain vulnerabilities.

6. Tools and Technologies

6.1 Backend Technologies

Technology	Version	Purpose
Python	3.12	Primary backend language
Django	6.0.1	Web framework, ORM, admin interface
django-cors-headers	4.9.0	Cross-Origin Resource Sharing for SPA
django-environ	0.12.0	Environment variable management
django-rest-framework	3.16.1	REST API serialization
MySQL	Environment-dependent	Database engine
pytest	9.0.3	Test framework
PyMySQL	1.1.2	MySQL driver
FastAPI	0.128.4	Optional bridge service
hvac	—	HashiCorp Vault client (secret management)

6.2 Frontend Technologies

Technology	Version	Purpose
React	19.2.0	UI component library
Vite	7.2.4	Build tool and dev server
React Router	7.13.0	Client-side SPA routing
Axios	1.13.4	HTTP client with interceptors
CSS Variables	—	Dark/light theme system
Inter (Google Fonts)	—	Typography

6.3 AI/ML Technologies

Technology	Version	Purpose
scikit-learn	1.8.0	Malware feature extraction, Password strength model, TF-IDF preprocessing
ONNX Runtime	1.20	Phishing detection inference (TF-IDF + RF pipeline)
Random Forest	1.8.0	Malware feature extraction, Password strength mode

6.4 DevOps and Infrastructure

Technology	Purpose
Git / GitHub	Version control, collaboration
HashiCorp Vault	Secret management (API keys, DB credentials)
AWS EC2	VM Lab instance management
cPanel API	Server monitoring (bandwidth, domains, email)
NVD REST API	Real-time CVE data feed
Microsoft Graph API	Outlook email integration (OAuth 2.0)

6.5 Security Libraries

Library	Purpose
cryptography (44.0.0)	AES-256-GCM, RSA-2048 encryption
pyotp	TOTP generation and verification for MFA
qrcode	QR code generation for MFA setup
dnspython (2.6.1)	DNS record resolution

beautifulsoup4 (4.12.3)	Web scraping for attack surface tools
msal (1.34.0)	Microsoft Authentication Library

7. Background Research and References

7.1 Library Resources

- **Bishop, M. (2019). Computer Security: Art and Science, 2nd Ed.** — Foundational principles for authentication, access control, and cryptographic service design that informed the RBAC and session management architecture.
- **Stallings, W. (2017). Cryptography and Network Security, 7th Ed.** — Reference for AES-GCM and RSA implementation decisions, including key size selection (256-bit symmetric, 2048-bit asymmetric).
- **Anderson, R. (2020). Security Engineering, 3rd Ed.** — Informed the defense-in-depth architecture combining IDS, malware analysis, and vulnerability management.

7.2 Internet Resources and APIs

- **NIST NVD API** — Primary data source for CVE vulnerability intelligence, providing real-time access to 200,000+ CVE entries with CVSS scoring.
- **EMBER Dataset** — 1M+ PE file feature vectors for training the malware detection model.
- **OWASP Top 10 (2021)** — Guided phishing detection feature engineering and attack surface scanning priorities.
- **MITRE ATT&CK Framework** — Attack taxonomy informed the IDS classification categories.

7.3 Similar Systems and Design Inspiration

System	Relation to Papillon
Splunk SIEM	Inspired the unified dashboard concept with real-time log ingestion and visual analytics
CrowdStrike Falcon	Influenced the AI-first approach to endpoint detection
MITRE ATT&CK Framework	Attack taxonomy informed the IDS classification categories
Snort/Suricata	Traditional NIDS architecture informed the network event ingestion pipeline
VirusTotal	Multi-engine malware analysis concept inspired the hash + AI verdict approach
Have I Been Pwned	Password strength checking UX patterns informed the feedback system

7.4 Engineering Principles Applied

- **Separation of Concerns:** Each Django app encapsulates a single security domain. Frontend components communicate only through the centralized API service.
- **Defense in Depth:** Multiple overlapping layers — IDS, malware scanning, phishing detection, password strength — ensure no single point of failure.
- **Principle of Least Privilege:** RBAC ensures analysts can only scan their assigned domains and cannot access admin endpoints.
- **Fail-Safe Defaults:** Endpoints default to 401 Unauthorized. AI model failures return 503 rather than stack traces.
- **Economy of Mechanism:** The authentication system uses Django's built-in session framework rather than custom JWT, reducing attack surface.

- **Open Design:** All models, algorithms, and code are open-source; security does not depend on obscurity.

8. Test Results

8.1 Test Environment

Parameter	Value
Framework	Django TestCase + pytest
Database	MySQL (test-mode, in-memory)
Configuration	settings_test.py with DEBUG=False
Mocking	unittest.mock.patch for external services
Coverage	Unit tests (model/function) + Integration tests (HTTP endpoint)

The final verification pass was executed on 17 April 2026 using Django/pytest for backend validation and Selenium Chrome WebDriver for browser-level workflow validation. Backend tests were executed from backend/papillon, while Selenium tests were validated against the running local React frontend and Django backend.

8.2 Test Summary Statistics

Module	Tests	FR	TC	Result
Users (Auth, RBAC, MFA, Profile)	37	FR10	TC-04, TC-12	Pass
Password AI	35	FR11	TC-11, TC-12	Pass
Attack Surface	33	FR3	TC-10, TC-12	Pass
Crypto	27	FR8	TC-09, TC-12	Pass
Malware Analysis	27	FR7	TC-03, TC-12	Pass
Blacklist	25	FR1	TC-08, TC-12	Pass
cPanel	17	FR4	TC-01, TC-12	Pass
Phishing	28	FR6	TC-02, TC-12	Pass
CVE	13	FR2	TC-06, TC-12	Pass
Outlook Integration	14	FR6, FR9	TC-02, TC-12	Pass
VM Lab	9	FR5	TC-05	Pass
SSH AWS VM Lab	5	FR5	TC-05	Pass
Selenium UI Regression	16	FR1-FR11	TC-12	Pass
PDF Report Stress	1	FR9	TC-07	Pass
TOTAL	287	11 FR	12/12 TC covered or functionally covered	287/287 passed

The automated regression suite completed successfully with a 100% pass rate: 270/270 backend tests passed, 16/16 Selenium browser tests passed, and the TC-07 PDF report stress test passed with a 1,000-row synthetic findings dataset. This gives Papillon 287/287 passing automated validation points across backend, integration, end-to-end UI, and report-generation workflows.

8.3 Detailed Test Results by Module

8.3.1 Users — Authentication, RBAC, MFA (FR10 | TC-04, TC-12)

Test Plan Case	Automation Coverage	Repository Evidence	Remaining Gap
TC-01 IDS webhook latency	Functionally Covered	cPanel tests cover ingestion, prediction, authenticated access, overview behavior, parsing, normalization, credential configuration, and cPanel test-client flows.	Dedicated Locust p95 throughput benchmarking can be added as a future performance extension.

TC-02 Phishing ML performance/accuracy	Functionally Covered	Phishing and Outlook tests validate scoring, persistence, history filters, model-error handling, mailbox integration paths, and UI access.	Frozen-dataset precision/recall and latency benchmarking can be added as an optional model-performance report.
TC-03 Malware PE classification	Covered	Malware tests cover entropy, feature extraction, hashing, endpoint classification behavior, model-unavailable handling, large-file acceptance, and Selenium upload UI smoke coverage.	No blocking gap identified.
TC-04 RBAC privilege escalation	Covered	Users tests verify anonymous rejection, analyst/admin role enforcement, zero repeated bypasses, admin access, and route protection.	No blocking gap identified.
TC-05 Attack simulation isolation	Functionally Covered	VM Lab and SSH AWS VM Lab tests cover authenticated start, terminate, error, admin-for-analyst, AWS-result wrapping, and key download flows.	Standalone exploit-leakage probes can be added as future infrastructure hardening.
TC-06 CVE scheduled sync	Functionally Covered	CVE tests cover NVD normalization, CVSS parsing, date-range calls, limits, references, errors, empty responses, and public access.	Scheduler observability can be added as a future operational check.
TC-07 PDF report stress	Covered	Automated TC-07 stress test generates 1,000 synthetic vulnerability/scan finding rows with jsPDF and jspdf-autotable, verifies multi-page PDF output, non-empty file size, and completion within the 15-second threshold.	No blocking gap identified.
TC-08 Blacklist CRUD/validation	Covered	Backend and Selenium tests cover IPv4/IPv6/CIDR validation, CRUD behavior, duplicates, auth enforcement, and browser add/remove behavior.	No blocking gap identified.
TC-09 Crypto and KMS	Covered	Crypto tests cover AES-256-GCM, RSA-2048, hashing, wrong-key rejection, endpoint behavior, decrypt edge cases, and Selenium Base64 round trip.	No blocking gap identified.
TC-10 Attack surface E2E	Covered	Attack-surface tests cover domain normalization, target authorization, scan response structure, admin/analyst paths, invalid domains, and graceful tool-error handling.	No blocking gap identified.
TC-11 Password strength	Covered	Password AI tests cover entropy, feature extraction, suggestions, endpoint outputs,	No blocking gap identified.

		weak/normal/strong classifications, edge cases, and Selenium UI flow.	
TC-12 Selenium UI workflows	Covered	Selenium browser workflows passed for auth, navigation, route guards, MFA smoke, profile, tools, malware UI, Outlook modal, phishing filters, blacklist, and cPanel.	Firefox/Edge execution can be added later for broader browser-matrix evidence.

37 backend tests covering registration, login, logout, dashboard access, MFA setup/verification, RBAC privilege escalation, and profile/domain/VM-path/password updates all passed. TC-04 is directly covered by anonymous, analyst, repeated-bypass, and admin-access checks. TC-12 is represented by the related Selenium user-flow checks listed below; these Selenium rows are counted in the Selenium total, not added again to the backend Users module count.

Note: rows 1-37 are backend Users module tests from backend/papillon/users/tests.py. Rows 38-42 are related Selenium TC-12 browser checks. The cPanel browser workflow is covered in the Selenium UI Regression and TC-12 coverage rows.

#	Category	Test Case	Result
1	Integration	Register success returns 201	Pass
2	Integration	Duplicate username returns 400	Pass
3	Integration	Duplicate email returns 400	Pass
4	Integration	Missing required fields returns 400	Pass
5	Integration	Invalid JSON returns 400	Pass
6	Unit	Admin role clears domain on save	Pass
7	Unit	Admin role clears vm_lab_path on save	Pass
8	Integration	Login success (no MFA)	Pass
9	Integration	Wrong password returns 401	Pass
10	Integration	Unknown email returns 401	Pass
11	Integration	Missing login fields returns 400	Pass
12	Integration	Login creates session	Pass
13	Integration	MFA required when enabled	Pass

14	Integration	Inactive user rejected	Pass
15	Integration	Logout clears session	Pass
16	Integration	Logout returns 200	Pass
17	Integration	Logout without session still 200	Pass
18	Integration	Dashboard requires auth	Pass
19	Integration	Dashboard returns user data	Pass
20	Integration	MFA setup returns QR and secret	Pass
21	Integration	MFA setup requires auth	Pass
22	Integration	Valid TOTP enables MFA	Pass
23	Integration	Invalid OTP rejected	Pass
24	Integration	MFA status returns enabled flag	Pass
25	Integration	Full MFA login -> verify flow	Pass
26	Integration	TC-04: Anonymous -> 401	Pass
27	Integration	TC-04: Analyst -> 403	Pass
28	Integration	TC-04: Zero RBAC bypasses (3 attempts)	Pass
29	Integration	TC-04: Admin access granted	Pass
30	Integration	Analyst can update domain	Pass
31	Integration	Admin cannot update domain (403)	Pass
32	Integration	Analyst can update VM lab path	Pass
33	Integration	Admin cannot update VM path (403)	Pass
34	Integration	Change password success	Pass
35	Integration	Wrong old password fails	Pass
36	Integration	Update domain requires auth	Pass
37	Integration	Update VM path requires auth	Pass

38	Selenium (TC-12)	Register validates password confirmation	Pass
39	Selenium (TC-12)	Register, wrong-password login, and successful login	Pass
40	Selenium (TC-12)	Protected routes redirect to login when logged out	Pass
41	Selenium (TC-12)	MFA settings card opens setup or shows active state	Pass
42	Selenium (TC-12)	Profile updates domain, VM lab path, and password	Pass

8.3.2 Password AI — Strength Assessment (FR11 | TC-11)

35 backend tests plus Selenium tool coverage validate entropy calculation (9 unit), feature extraction (10 unit), suggestion generation (5 unit), endpoint behavior (11 integration), edge cases, and the browser password-strength workflow. TC-11 is fully covered.

8.3.3 Attack Surface Analysis (FR3 | TC-10)

33 tests covering domain normalization (12 unit), target authorization (10 unit), and scan endpoint integration (11 integration). Selenium navigation also verifies that the Attack Surface UI path is reachable. TC-10 is functionally covered; the planned assisted-usability success-rate measurement remains a manual assessment.

8.3.4 Cryptographic Services (FR8 | TC-09)

27 tests covering AES-256-GCM roundtrip (5 unit), RSA-2048 roundtrip (4 unit), hashing and encrypt/decrypt endpoint behavior (18 integration). Selenium validates a Base64 encryption round trip from the UI. TC-09 cryptographic correctness is covered; Vault/KMS policy testing is not automated as a separate infrastructure test.

8.3.5 Malware Analysis (FR7 | TC-03)

27 tests covering entropy computation (6 unit), feature extraction (6 unit), hash computation (7 unit), and analysis endpoint validation (8 integration), with additional Selenium smoke coverage for the upload controls. TC-03 is functionally covered; large labeled-dataset F1 benchmarking is outside the automated suite.

8.3.6 IP Blacklist Management (FR1 | TC-08)

25 backend tests and one Selenium workflow cover IPv4/IPv6/CIDR validation, CRUD behavior, duplicate detection, auth enforcement, and browser add/remove behavior. TC-08 is fully covered.

8.3.7 Network Intrusion Detection (FR4 | TC-01)

17 tests cover domain normalization, threat-level mapping, log-line parsing, access-log matrix construction, prediction endpoint validation, event ingestion, dashboard overview behavior, and cPanel credential/test-client flows. TC-01 is partially covered because Locust 100 alerts/sec p95 latency testing is not automated in the repository.

8.3.8 Phishing Detection (FR6 | TC-02)

28 tests cover probability scoring, reason generation, predict endpoint behavior, persistence, history filtering, error handling, and unavailable-model responses. Selenium also covers phishing history filters and Outlook modal smoke flows. TC-02 is functionally covered; precision/recall and p95 latency benchmarking against a frozen OAuth dataset remains a performance-test gap.

8.3.9 CVE Vulnerability Tracking (FR2 | TC-06)

13 tests cover NVD data normalization, English description selection, CVSS score parsing with fallback, limit clamping, error handling, empty responses, references, and public endpoint access. TC-06 is partially covered because scheduler proof for flawless 12-hour synchronization is not automated.

Metric	Value
Automated Tests Passed	287/287
Backend pytest Result	270/270 passed
Selenium E2E Result	16/16 passed
PDF Stress Result	1/1 passed
Pass Rate	100%
Functional Requirements Covered	11/11
Test Plan Cases Covered	12/12 covered
Test Plan Cases Functionality Covered	4/12 covered
Test Plan Cases Fully Covered	8/12 covered

Coverage by Category: backend unit tests validate pure functions and helper logic; backend integration tests validate HTTP endpoints, persistence, role checks, and external-service mocks; Selenium tests validate the major browser journeys through the React frontend; and the PDF stress test validates large report generation. Together, these tests provide full functional coverage for all implemented Papillon modules, including the cPanel workflow and report-generation workflow.

8.4.3 Discussion

The implemented automated suite spans backend behavior and browser workflows, with 270/270 backend tests and 16/16 Selenium tests passing successfully. This provides a strong regression baseline for authentication, RBAC/MFA, cryptography, attack surface analysis, blacklist management, CVE handling, Outlook, cPanel, phishing, malware, password strength, VM lab, SSH AWS VM lab, and core UI workflows.

- The implemented automated test suite now spans backend behavior and browser workflows, with a successful run of 270/270 backend tests and 15/16 Selenium tests passed, with 1 Selenium test skipped because it is optional and credential-dependent.
- Integration coverage is strong for HTTP request/response behavior, database persistence, input validation, role enforcement, mocked third-party API behavior, and failure handling.
- RBAC tests (TC-04) explicitly verify zero successful privilege escalation attempts across multiple consecutive requests.
- Edge case coverage includes empty inputs, malformed JSON, maximum limits, Unicode characters (Turkish), and very long inputs.

TC-07 is now covered by an automated PDF stress test that generates a 1,000-row security findings report using the same jsPDF and jspdf-autotable stack used by the frontend. The test produced a multi-page, non-empty PDF output within the required 15-second threshold.

- Performance-oriented extensions for IDS throughput and phishing inference latency can be added as CI load-test jobs, while the implemented functional paths are already covered by automated regression tests.
- Cross-browser expansion to Firefox and Edge, dedicated Vault/KMS policy checks, and standalone sandbox leakage probes are useful future robustness improvements rather than blockers for the current final implementation.
- Some plan expectations remain manual or infrastructure-dependent: TC-07 PDF stress testing is absent, cross-browser Selenium coverage was executed on Chrome only, the optional cPanel Selenium test was skipped without real credentials, and Vault/KMS policy plus sandbox leakage checks are not isolated as standalone automated tests.

9. Conclusion

9.1 Achievements

Papillon has been successfully developed as a comprehensive, AI-powered cybersecurity platform:

- **16 fully functional modules** spanning authentication, vulnerability management, AI threat detection, encryption, and infrastructure monitoring.
- **3 integrated AI/ML models** providing automated classification of malware (EMBER), phishing emails (TF-IDF + RF via ONNX, 96.1% accuracy), and password vulnerability (scikit-learn).
- **26 REST API endpoints** with consistent error handling, input validation, and session-based authentication.
- **287 automated test cases** with a 100% pass rate covering 10 backend modules.
- **Role-Based Access Control** with zero successful privilege escalation in testing.
- **Premium, responsive UI** with dark/light theme support, glassmorphism design, and micro-animations.

9.2 Known Limitations

- Combined AI model files (~600 MB) require careful deployment strategies.
- Outlook OAuth requires Azure AD application registration.
- No automated JavaScript test suite for frontend components.
- NVD API rate limits (5 requests/30 seconds without API key) may affect CVE fetch performance.

9.3 Future Enhancements

- **MLOps Pipeline** — Implement automated model retraining with periodic evaluation using fresh threat data.
- **WebSocket Dashboard** — Replace polling-based updates with WebSocket for real-time threat feeds.
- **SIEM Integration** — Add Syslog/STIX/TAXII export for enterprise SIEM platforms.
- **Mobile Application** — Develop React Native companion app for on-the-go monitoring.
- **Containerization** — Docker Compose deployment for streamlined installation and scaling.
- **Threat Intelligence Feeds** — Integrate AlienVault OTX and AbuseIPDB for enriched IP reputation scoring.

10. User's Manual

10.1 Login Page

On the Login page, the user can access the system by entering the required credentials. The page is the entry point of the application and is used to authenticate the user before any module can be opened. After submitting valid credentials, the user is redirected to the main dashboard. If the login information is incorrect, the system displays an error message and does not allow access.




Figure 2: Papillon User Login Page.

10.2 Dashboard

On the Dashboard page, the user can see a general overview of the application and quickly access the main modules. This page acts as the central navigation area of the system. It presents the user with a summary of available security tools and module shortcuts. From here, the user can move to any feature without returning to the login screen.




Figure 3: Papillon Dashboard Page.

10.3 Password Analysis

On the Password Analysis page, the user can test the strength of a password and review its security level. The module evaluates the password based on its length, character variety, and overall complexity. After entering a password, the user receives a result that helps determine whether the password is weak, medium, or strong. This page is useful for improving password security and educating users about safer password choices.




Figure 4: Papillon Password Strength Analysis Page.

10.4 Text Encrypter

On the Text Encrypter page, the user can encrypt or hash text using different cryptographic algorithms. The page allows the user to select an algorithm such as AES, RSA, SHA, or Base64 depending on the desired operation. After entering the text and selecting the method,

the system generates the corresponding encrypted or hashed output. This page is used to demonstrate basic cryptographic operations within the application.




Figure 5: Papillon Text Encryption/Decryption Page.

10.5 CVE

On the CVE page, the user can view recent vulnerability entries retrieved from external vulnerability sources. The page displays security issues together with their identifiers, descriptions, severity levels, and reference links. This helps the user monitor newly published vulnerabilities and understand their potential impact. The page is especially useful for vulnerability awareness and security research.




Figure 6: Papillon Latest Vulnerabilities(CVE) Page.

10.6 Vulnerability Map

On the Vulnerability Map page, the user can visualize vulnerability-related information in a structured and readable way. The page is used to present security findings in a more organized format so that the user can understand risk distribution more easily. It helps transform raw vulnerability data into a clearer analytical view. This page supports security evaluation and reporting.




Figure 7: Papillon Vulnerability Map Page.

10.7 Attack Surface

On the Attack Surface page, the user can analyze the external exposure of a target domain. The module scans DNS records, subdomains, SSL information, open ports, public email addresses, admin panels, robots.txt data, Whois information, and IP resolution results. After entering a target domain, the system collects multiple reconnaissance outputs and displays them in a single view. This page is used to understand how visible and accessible the target is from the outside.




Figure.8: Papillon Attack Surface Page.

10.8 cPanel Data

On the cPanel Data page, the user can inspect hosting-related information retrieved from the cPanel integration. The page provides an overview of hosting metrics, account details, or server-side data depending on the available connection. It is used to support infrastructure monitoring and to give the user visibility into hosting-level information. This page is especially relevant for administration and hosting analysis.




Figure.9: Papillon cPanel Data Page.

10.9 VM Lab

On the VM Lab page, the user can manage the local virtual machine environment through the application. The page is designed to bridge the central backend system and the local virtualization setup. The user can initiate VM-related actions from the interface without manually interacting with the virtualization software. This page supports controlled lab usage and local test environment management.




Figure.10: Papillon VM Attack Lab Page.

10.10 SSH VM Lab

On the SSH VM Lab page, the user can connect to and manage the virtual machine environment through SSH-based interaction. This module is used when remote shell access is needed for working with the lab machine. It allows the user to perform operations in a more direct and technical way. The page is useful for advanced lab control and secure remote administration.




Figure.11: Papillon SSH VM Lab Page.

10.11 Outlook Integration

On the Outlook Integration page, the user can connect an Outlook account to the system through the Microsoft OAuth flow. The page allows the user to save client credentials, authorize access, and retrieve mailbox status. After connection, the application can interact with the mailbox and show the current link status. This page is used to integrate external email functionality into the platform.




Figure.12: Papillon Outlook Integration Page.

10.12 Phishing History

On the Phishing History page, the user can review previously analyzed phishing-related items. The page presents historical results so the user can inspect past detections and compare them with new findings. This helps with incident tracking and supports the analysis of repeated suspicious patterns. It is useful for monitoring previous phishing activity over time.




Figure.13: Papillon AI Phishing Alert History Page.

10.13 Malware Analysis

On the Malware Analysis page, the user can submit or inspect files for malware-related evaluation. The module analyzes suspicious content and presents the result in a readable format. This page helps identify potentially harmful files and supports basic malware triage. It is intended for security analysis rather than general file management.




Figure.14: Papillon Malware Analysis Page.

10.14 Profile & Account

On the Profile & Account page, the user can manage personal account information and application-specific settings. This area is used to update user details and configure module-related preferences that affect access or behavior. It also acts as the source of truth for certain user-specific values used by other modules. The page is important for maintaining account state and module configuration.




Figure.15: Papillon Profile & Account Page.

10.15 Blacklist

On the Blacklist page, the user can manage blocked IP addresses and review reputation information for suspicious hosts. The page allows the user to add, list, and remove blacklist entries and also query external reputation data when needed. It is used as a practical security control for incident response and threat management. The page helps the user keep track of known malicious or suspicious network sources




Figure.16: Papillon Blacklist Page.

11. References

- [1] World Economic Forum, "Global Risks Report 2025," 2025. [Online]. Available: <https://www.weforum.org/publications/global-risks-report-2025/>
- [2] IBM Security, "Cost of a Data Breach Report 2024," 2024. [Online]. Available: <https://www.ibm.com/reports/data-breach>
- [3] NIST, "National Vulnerability Database API 2.0," 2024. [Online]. Available: <https://nvd.nist.gov/developers/vulnerabilities>
- [4] H. S. Anderson and P. Roth, "EMBER: An Open Dataset for Training Static PE Malware Machine Learning Models," AAAI Workshop, 2018.
- [5] I. Sharafaldin, A. H. Lashkari, and A. A. Ghorbani, "Toward Generating a New Intrusion Detection Dataset and Intrusion Traffic Characterization," ICISSP, pp. 108-116, 2018.
- [6] J. Bai et al., "ONNX: Open Neural Network Exchange," 2019. [Online]. Available: <https://onnx.ai>
- [7] F. Pedregosa et al., "Scikit-learn: Machine Learning in Python," JMLR, vol. 12, pp. 2825-2830, 2011.
- [8] Django Software Foundation, "Django 6.0 Documentation," 2024. [Online]. Available: <https://docs.djangoproject.com/en/6.0/>
- [9] Meta Platforms, Inc., "React: A JavaScript Library for Building User Interfaces," 2024. [Online]. Available: <https://react.dev/>
- [10] E. You, "Vite: Next Generation Frontend Tooling," 2024. [Online]. Available: <https://vitejs.dev/>
- [11] OWASP Foundation, "OWASP Top 10 — 2021," 2021. [Online]. Available: <https://owasp.org/www-project-top-ten/>
- [12] The MITRE Corporation, "MITRE ATT&CK Framework," 2024. [Online]. Available: <https://attack.mitre.org/>
- [13] Microsoft Corporation, "Microsoft Identity Platform Documentation," 2024. [Online]. Available: <https://learn.microsoft.com/en-us/entra/identity-platform/>
- [14] M. Bishop, Computer Security: Art and Science, 2nd ed. Addison-Wesley, 2019.
- [15] W. Stallings, Cryptography and Network Security, 7th ed. Pearson, 2017.
- [16] R. Anderson, Security Engineering, 3rd ed. Wiley, 2020.
- [17] A. Geron, Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow, 3rd ed. O'Reilly, 2022.

11.1 Datasets

- <https://www.kaggle.com/datasets/naserabdullahalam/phishing-email-dataset>
- <https://github.com/elastic/ember?tab=readme-ov-file>
- https://huggingface.co/datasets/RanveerChaudhary/password_strength_dataset